

## **Artificial Neural Networks Programming Assignment 2**

**DUE DATE: DEC 30, 2019**

**OBJECT OF THE ASSIGNMENT:**

To understand how the **Backpropagation algorithm** learns the weight and bias values for *multilayer networks*. And to realize how different numbers of hidden nodes and different learning rates change the performance of the backpropagation algorithm.

**PROBLEM:**

Implement the **Stochastic Backpropagation** to classify the popular UCI Iris dataset. Note that you should implement the algorithm from scratch without using any other existing software.

**DATASETS:**

Your TA will split the Iris dataset into two sub-datasets, 120 examples of the training dataset and 30 examples of the testing dataset. Both datasets have the following characteristics:

- Four numeric attributes: sepal length, sepal width, petal length, and petal width
- Three class labels: Iris-setosa, Iris-versicolor, and Iris-virginica

**INPUT OF THE PROBLEM:**

Training dataset/testing dataset

**OUTPUT OF THE PROBLEM:**

Display both of the training and testing accuracies and the number of epochs when the program stops.

**EXPERIMENTS:**

- (a) Build a two-layer neural network with four components in the input layer and three neurons in the output layer. Find experimentally a good number of hidden neurons, start with the number of hidden neurons = 1.
- (b) Rerun the experiment for different learning rates.

## DISCUSSION:

Write a brief report:

- (a) Report which neural network architecture (i.e., different number of hidden nodes) obtained the best performance. Analyze and explain your findings.
- (b) Compare the performance of different learning rates.

You may summarize your experimental results using a table or a figure to discuss and analyze your results.

## EXTRA CREDIT PORTION (+20 POINTS)

Derive the gradient descent rule for the Softmax activation using the cross-entropy loss and show the derivation in the report. Implement the **Stochastic Backpropagation** and rerun the experiments. Is the performance better than the Sigmoid activation using the squared-error loss? Discuss and explain your findings.

## REMARKS

1. Stopping criteria usually includes:

- (a) Stop when a maximum number of epochs has been exceeded.
- (b) Stop when the root mean squared error (RMSE) on the training set  $D$  is small enough (other error measures such as the mean absolute error (MAE) can also be used).

$$- E_{RMSE} = \sum_{d \in D} \sqrt{\frac{(t_d - o_d)^2}{|D|}} < \text{a given } \tau_1$$

$$- E_{MAE} = \sum_{d \in D} |t_d - o_d| / |D| < \text{a given } \tau_2$$

where  $|D|$ : the number of training examples

$t_d$ : the target of the training example  $d$

$o_d$ : the neuron output of the training example  $d$

2. One obvious choice of the target vectors for three class labels are:  $[1 \ 0 \ 0]^T$ ,  $[0 \ 1 \ 0]^T$ , and  $[0 \ 0 \ 1]^T$ . Instead of 0 and 1 values, use values of 0.1 and 0.9, so that the target vectors are  $[0.9 \ 0.1 \ 0.1]^T$ ,  $[0.1 \ 0.9 \ 0.1]^T$ , and  $[0.1 \ 0.1 \ 0.9]^T$ . The reason for avoiding target values of 0 and 1 is that sigmoid units cannot produce these output values given finite weights.